

COSC 511: Computer Architecture

Arithmetic for Computers

Week 3

Last Week

- Instruction Set Architectures (ISAs)
 - Examples of ISAs (x86, ARM)
 - Tradeoffs between different types of ISAs
 - Compatibility between ISAs
 - Real-time ISA translation use cases, benefits, and limitations
- ISA Design Principles
- ISAs optimize for the computer, not for the human.
 - Example: Twos complement for handling negative integers.
- High-Level Language Optimization
 - Example: **gcc** allows programmers to control code optimization.
- Interesting side effects of ISA implementations
 - Illegal Opcodes

Arithmetic for Computers

- The obvious stuff:
 - Addition and Subtraction
 - Multiplication and Division
 - Following Order of Operations
- The not-so-obvious stuff:
 - Dealing with overflow
 - Representing fractional (floating point) values in binary

Arithmetic for Computers – Addition

- Remember:

- $0 + 0 = 0$

- $0 + 1 = 1$

- $1 + 1 = 0$, carry 1

- $1 + 1 + 1 = 1$, carry 1

$$\begin{array}{r} 1 \quad 1111 \\ 01010010 \quad 82 \\ + 01001110 \quad 78 \\ \hline 10100000 \quad 160 \end{array}$$

Arithmetic for Computers – Subtraction

- Remember:

$$- 0 - 0 = 0$$

$$- 1 - 0 = 1$$

$$- 1 - 1 = 0$$

$$- 0 - 1 = 1, \text{ borrow } 1$$

This works, but it's not how computers do it.

$$\begin{array}{r} 1 \\ 010\cancel{0}\cancel{0}010 \quad 82 \\ - 00001110 \quad 14 \\ \hline 01000100 \quad 68 \end{array}$$

Arithmetic for Computers – Subtraction

- Remember: Make ISA design as simple as possible.
 - Is it easier to implement subtraction separately at the hardware level?
 - No, it is not.
 - It is much better to use a cool little trick to use addition to do subtraction.
 - Twos compliment!

Arithmetic for Computers – Subtraction

- Twos Complement
 1. Apply bitwise NOT to the integer representation of the number.
 2. Add one to it.
 3. Add values together.

00001110

Arithmetic for Computers – Subtraction

- Twos Complement
 1. Apply bitwise NOT to the integer representation of the number.
 2. Add one to it.
 3. Add values together.

00001110
11111111

Arithmetic for Computers – Subtraction

- Twos Complement
 1. Apply bitwise NOT to the integer representation of the number.
 2. Add one to it.
 3. Add values together.

$$\begin{array}{r} 1 \\ 11110001 \\ + 1 \\ \hline 11110010 \end{array}$$

Arithmetic for Computers – Subtraction

- Remember:

$$- 0 + 0 = 0$$

$$- 0 + 1 = 1$$

$$- 1 + 1 = 0, \text{ carry } 1$$

$$- 1 + 1 + 1 = 1, \text{ carry } 1$$

$$\begin{array}{r} 1 \ 111 \ 1 \\ 01010010 \ 82 \\ + 11110010 \ -14 \\ \hline 01000100 \ 68 \end{array}$$

Arithmetic for Computers – Multiplication

- Remember:

- $0 \times 0 = 0$

- $0 \times 1 = 0$

- $1 \times 0 = 0$

- $1 \times 1 = 1$

$$\begin{array}{r} 11101 \\ \times 01001 \\ \hline \end{array}$$

Arithmetic for Computers – Multiplication

- Remember:

- $0 \times 0 = 0$

- $0 \times 1 = 0$

- $1 \times 0 = 0$

- $1 \times 1 = 1$

$$\begin{array}{r} 11101 \\ \times 01001 \\ \hline 11101 \\ 000 \end{array}$$

Arithmetic for Computers – Multiplication

- Remember:

- $0 \times 0 = 0$

- $0 \times 1 = 0$

- $1 \times 0 = 0$

- $1 \times 1 = 1$

$$\begin{array}{r} 11101 \\ \times 01001 \\ \hline 11101 \\ + 11101000 \\ \hline \end{array}$$

Arithmetic for Computers – Multiplication

- Remember:

$$- 0 \times 0 = 0$$

$$- 0 \times 1 = 0$$

$$- 1 \times 0 = 0$$

$$- 1 \times 1 = 1$$

$$\begin{array}{r} 11101 \\ + 11101000 \\ \hline 100000101 \end{array}$$

How do computers make it easy to do this?

Answer: Shifting, then adding!

Arithmetic for Computers – Multiplication

- Remember:

- $0 \times 0 = 0$

- $0 \times 1 = 0$

- $1 \times 0 = 0$

- $1 \times 1 = 1$

$$\begin{array}{r} 11101 \\ \times 01001 \\ \hline 11101 \\ + 11101000 \\ \hline \end{array}$$

Shift left 3 spaces.

Arithmetic for Computers – Division

- Division is hard for computers!
 - Programmers who work on building systems where tiny fractions of time matter try to avoid using division because it is slower than other operations.
 - There are two many division implementations: slow and fast division
- Slow Division
 - Keep subtracting until you get to a point at which you cannot subtract anymore.
 - Example: $7 \div 3$
 1. $7 - 3 = 4$
 2. $4 - 3 = 1$
 3. $1 < 3$
 4. So, $7 \div 3 == 2$ and $7 \% 3 == 1$

Arithmetic for Computers – Division

- Fast Division
 - There are many different fast division methods.
 - These fast division methods take less computation time but rely on estimating the final value.
 - Example: Goldschmidt Fast Division
 - AMD has used this approach in their CPUs since the launch of AMD Athlon in June 1999.

Arithmetic for Computers

- Order of Operations
 - Computers follow the same approach to order of operations that we use.
- How do computers handle order of operations?
 - They don't! That's your problem.
 - At least, it is if you're writing assembly.
- Compilers are responsible for generating Assembly code that obeys order of operations.

$$(5 + 3) / 45 * 3 \% 9 - 2 / 52$$

Arithmetic for Computers

Assume we are using a high-level language, and the datatype is byte.

$$\begin{array}{r} 1111111 \\ +0000001 \\ \hline 10000000 \end{array}$$

What's wrong?

We can't fit this in a byte. ☹️

Arithmetic for Computers

- Year 2038 Problem
 - Many systems store time using the Unix epoch
 - Unix epoch: An integer value representing the number of seconds that have passed since January 1, 1970 at 00:00 UTC
 - At the start of today's class, the Unix epoch was 1694037600
 - Many systems use signed 32-bit integers for storing the Unix epoch
 - One integer is lost due to signing, so 31 bits are available for storing the timestamp.
 - On January 19, 2038 at 03:14:07 UTC, the timestamp will be:
 - **01111111 11111111 11111111 11111111**
 - On January 19, 2038 at 03:14:08 UTC, the timestamp will be:
 - **10000000 00000000 00000000 00000000**
 - Overflow!

Arithmetic for Computers

- Year 2038 Problem

- Overflow!
- Systems that use signed integers will travel back in time.
 - They will report time as December 13, 1901 at 20:45:52 UTC

```
Binary   : 01111111 11111111 11111111 11110000
Decimal  : 2147483632
Date     : 2038-01-19 03:13:52 (UTC)
Date     : 2038-01-19 03:13:52 (UTC)
```

- Solutions

- Use unsigned 32-bit integer
 - We'd be fine until February 7, 2106 at 06:28:15 UTC
- Use a 64-bit integer
 - We'll be fine for approximately 292 billion years.

Arithmetic for Computers

- Back to our addition problem: How do we fix this?
 - Use a short instead of a byte? (2 bytes)
 - Use an int? (4 bytes)
 - Use a long? (8 bytes)
 - But what if that's still not enough?
- That's what Binary Coded Decimal (BCD) is for!
 - BCD stores each digit in a value separately.
 - Java's **BigInteger** is a BCD implementation.

Arithmetic for Computers

- BCD of 123
 - Big-Endian Order – Leftmost digit comes first.
 - 00000001 00000010 00000011
 - Little-Endian Order – Rightmost digit comes first.
 - 00000011 00000010 00000001
- This is great, but it wastes space. Why?
 - We only ever use the first four bits of our byte!
 - This is called unpacked BCD.

0 = 0000

1 = 0001

2 = 0010

3 = 0011

4 = 0100

5 = 0101

6 = 0110

7 = 0111

8 = 1000

9 = 1001

Arithmetic for Computers

- Packed BCD
 - Packed BCD is more efficient.
 - Use one byte to store two digits.
- Packed BCD of 123
 - Big-Endian Order – Leftmost digit comes first.
 - **00000001 00100011**
 - Little-Endian Order – Rightmost digit comes first.
 - **00110010 00010000**

Arithmetic for Computers

- Pros of BCD
 - You can store really large values!
 - BCD is useful for dealing with rounding errors.
- Cons of BCD
 - Harder to do math operations with.
 - Depending on the value you are storing, more space is used.
 - Example: **123**
 - Without BCD: **01111011**
 - Big-Endian Packed BCD: **00000001 00100011**

Arithmetic for Computers

- Your data is what you make of it.
 - 00000001 00100011 has meaning as big-endian packed BCD.
 - 00000001 00100011 has meaning as a short.
 - This is why remote code execution vulnerabilities are a problem.
 - The same data can be interpreted in different ways.

```
public class ArrayProblem {  
    public static void main(String[] args) {  
        int ar[] = { 1, 2, 3, 4, 5 };  
        for (int i = 0; i <= ar.length; i++)  
            System.out.println(ar[i]);  
    }  
}
```

Arithmetic for Computers

- Your data is what you make of it.
 - 00000001 00100011 has meaning as big-endian packed BCD.
 - 00000001 00100011 has meaning as a short.
 - This is why remote code execution vulnerabilities are a problem.
 - The same data can be interpreted in different ways.
- If you write code that writes other data stored in memory adjacent to the intended integer value. The danger is in writing. When I store data in memory, I can deliberately exceed the bounds of the data stored in memory to be executed. My program will malfunction and/or crash. I can write code into an otherwise safe application.

```
#include <iostream>
using namespace std;

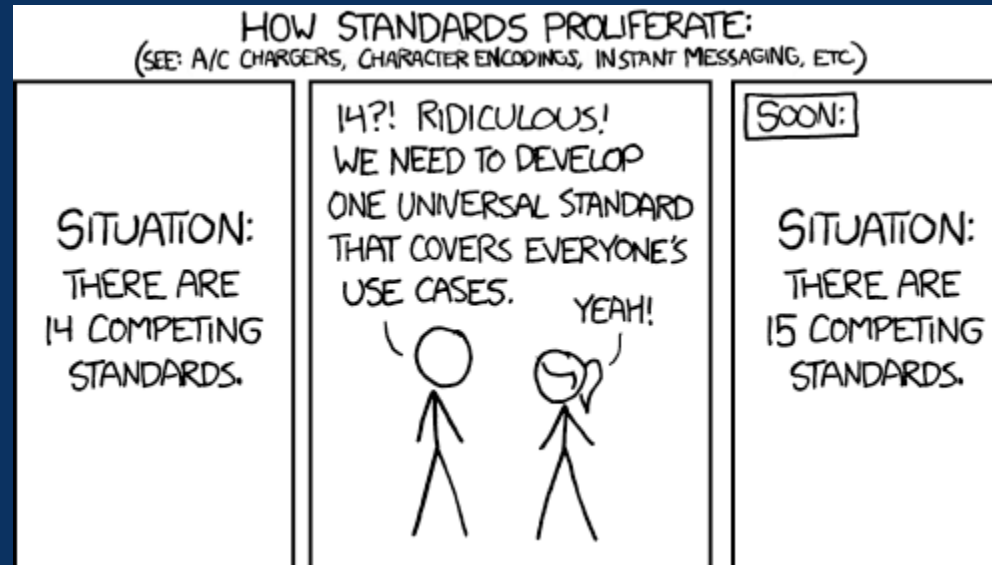
int main()
{
    int arr[] = {1,2,3,4,5};
    for (int x = 0; x <= 10; x++)
        cout << arr[x] << endl;

    return 0;
}
```

```
0
2127166208 452279552
-137567531 31503784
0
-1557131888 -987005552
```

Arithmetic for Computers

- Storing floating point values in binary is very confusing.
 - Many different implementations for storing floating point values used to exist.



- IEEE Standard for Floating-Point Arithmetic (IEEE 754), established in 1985, has become the standard implementation.

Arithmetic for Computers

- IEEE 754 defines two sizes for storing floating point values
 - Single precision numbers
 - Use 32 bits
 - We usually call these “floats”
 - Double precision numbers
 - Use 64 bits
 - We usually call these “doubles”
- IEEE 754 Floating Point Components
 1. Sign bit (0 for positive, 1 for negative)
 2. Exponent
 - 8 bits used for single, 11 bits for double
 3. Mantissa
 - 23 bits used for single, 52 bits for double

Arithmetic for Computers

- In most cases today, doubles are used for better precision.
- Converting to floating point values to binary:
 1. Convert the whole number portion to binary as you normally would.

85.125

85 == 1010101₂

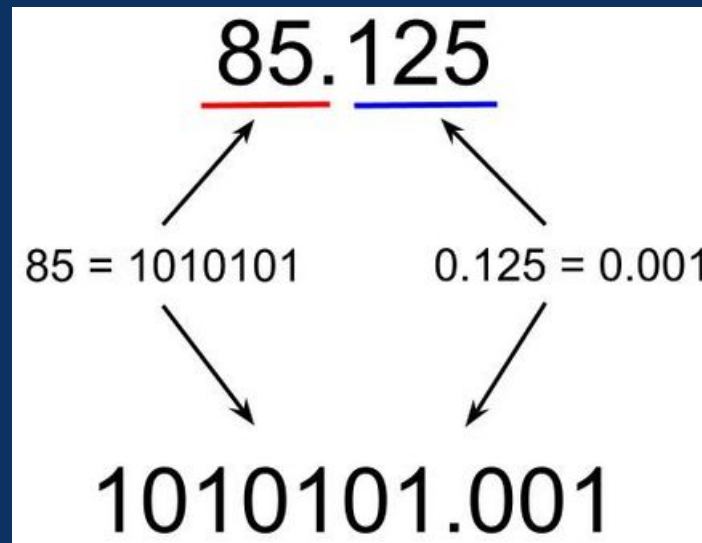
2. Convert the decimal portion to binary.

Decimal Number Multiplication	Result	Number in front of decimal
0.125 x 2	0.25	0
0.25 x 2	0.5	0
0.5 x 2	1.0	1
0.0 x 2	0.0	0

0.125 = 001

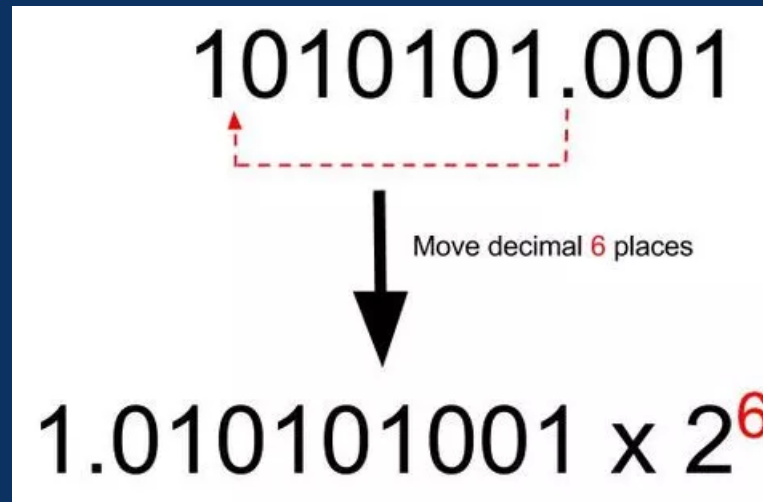
Arithmetic for Computers

- Converting to floating point values to binary:
 3. For organizational purposes, concatenate both binary components together.



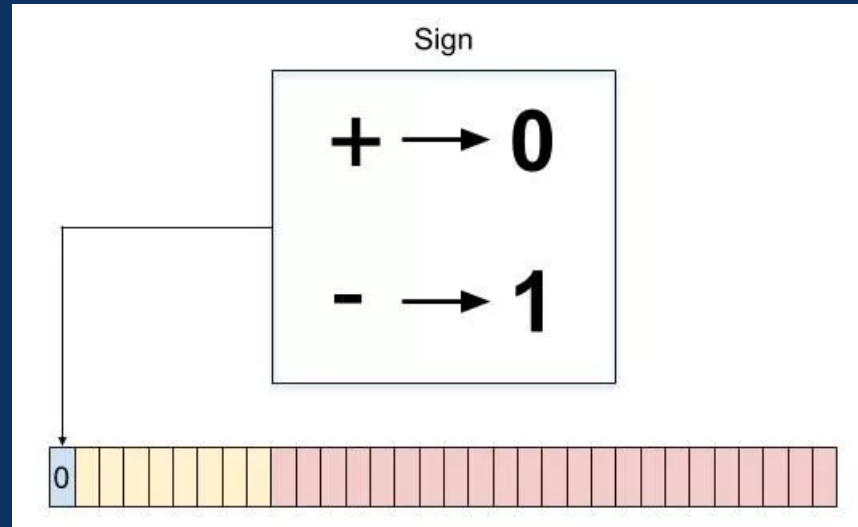
Arithmetic for Computers

- Converting to floating point values to binary:
 4. Convert the binary value to base 2 scientific notation.



Arithmetic for Computers

- Converting to floating point values to binary:
 5. Based on the sign of the original number, take note of the most significant bit.



Arithmetic for Computers

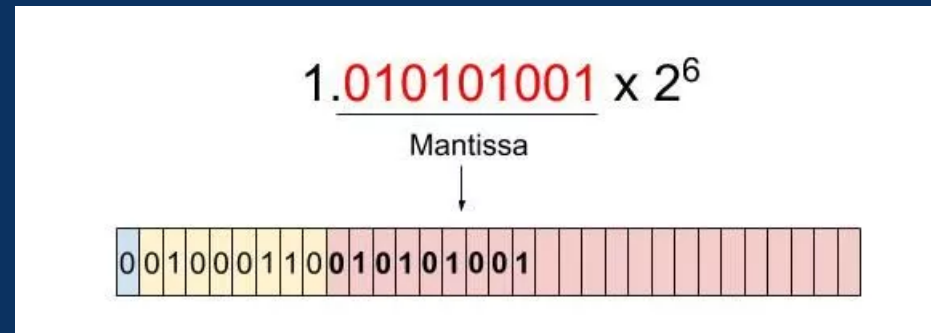
- Converting to floating point values to binary:
 6. Determine the exponent to use by applying the “exponent bias”
 - The purpose of exponent bias is to prevent issues with floating point signing
 - If you are using single precision, you calculate bias by adding 127 to the exponent from step 4.
 - If you are using double precision, you calculate bias by adding 1023 to the exponent from step 4.

$$1.010101001 \times 2^6$$

$$127 + 6 = 133$$

Arithmetic for Computers

- Converting to floating point values to binary:
 7. Determine the mantissa
 - Mantissa – Refers to the binary component after the decimal point in your scientific notation representation.



Arithmetic for Computers

- Converting to floating point values to binary:
 8. Assemble the components of your calculations together!

