

COSC 511: Computer Architecture

The Processor (Part 2)

Week 5

Last Week

- The Processor (Part 1)
 - Generic architecture of CPUs
 - The CPU execution cycle
 - Improving performance by pipelining the CPU execution cycle
 - Use of caching at multiple levels to improve performance
 - Creating “faster” CPUs by adding more cores
 - Understanding tradeoffs between higher clock speeds and more cores
 - Simultaneous Multithreading
 - Branch Prediction
 - Spectre: A security vulnerability in most CPUs made before 2019 caused by how branch prediction was implemented

The Processor

- Loading and Executing a Program
 - Programs are stored on the system's Hard Drive/Solid State Drive
 - To load and execute a program:
 1. OS tells the hard drive to locate and retrieve the application
 2. OS finds free RAM and copies the application to it
 3. OS creates a process that has the job of executing the program starting at its entry point
 - The entry point is the memory address where the program begins
 4. The process begins executing the application at its entry point
 - The process will take care of running the program
 - The OS has the job of tracking the execution process and responding to requests for system resources (memory allocation, file I/O, etc.)
 5. When the process ends, it is removed from memory

The Processor

- Allowing applications to have unrestricted access to memory can be a security risk.
 - Example:
 - My web browser is using memory between addresses **0x500000** and **0x700000**.
 - My e-mail client is using memory between addresses **0x700000** and **0x900000**.
 - I accidentally install malware on my computer that is loaded in memory between **0x910000** and **0x950000**.
 - Malware accesses memory allocated to my web browser and e-mail client.
 - With this access, it can:
 - Read private data from other applications' memory
 - Deliberately cause other applications to crash
 - Inject code into other applications' memory, causing these applications to execute other malicious code on its behalf
- Some ISAs implement different memory access modes to prevent this.

The Processor

- x86/amd64 Memory Access Modes
 - Real Address Mode
 - This is the mode described in the prior example.
 - Applications running in real address mode have direct access to all memory.
 - Memory addresses that a program “sees” correspond to real, physical memory locations.
 - Protected Mode
 - In protected mode, applications are not allowed to access memory that is not allocated to them.
 - Programs executed in protected mode are allocated a section of linear address space in memory and cannot access anything outside of it.
 - Memory addresses that a program “sees” are virtual addresses that are mapped to physical memory addresses under the hood.
 - Segmentation Faults!

The Processor

- Q: If real address mode is less secure, why do we still have it?
 - A: Some applications need it!
 - Examples:
 - Operating System components
 - Hardware Drivers
 - Antivirus Software
- Applications running in real address mode can switch to protected.
 - Switching from protected to real address mode is not possible.

The Processor

- Real Address Mode Pros
 - Easier to implement than protected mode
 - No additional memory management overhead
- Real Address Mode Cons
 - Huge security risk
 - Does not allow for any memory management

The Processor

- Protected Mode Pros
 - Allows for extensive control over memory management
 - Prevents deliberate or accidental tampering with other applications' memory
 - Exception: Malware that targets applications running in real address mode can circumvent this.
- Protected Mode Cons
 - Harder to implement than real address mode
 - Added complexity due to need for memory management

The Processor

- Memory Management
 - Maps real memory addresses to virtual memory addresses
 - Paging
 - Computer uses secondary storage to offload memory contents and reload them as needed
 - Paged data is either written to a page file or a dedicated drive partition
 - This is useful for when the system needs more memory than it has
 - Problem: Although paging is useful, it causes a noticeable performance bottleneck
 - The secondary storage is slower than RAM, so performance will be degraded
 - Performance degradation is less bad for computers with SSDs, but still present
 - For mechanical hard drives, performance degradation is worse because of higher seek times
 - » HDD average seek time: 9ms
 - » SSD average seek time: 0.16ms

The Processor

- Another Paging Problem
 - Paging causes lots of read and write operations to your hard drive
 - Old hard drives allowed for an infinite number of write cycles
 - Modern SSDs have a finite number of write cycles
 - An SSD that stores a single data bit per cell typically supports up to 100,000 write cycles.
 - Paging is a very write-heavy operation
 - Frequent paging can kill your SSD prematurely!

The Processor

- Paging: The SSD Killer – Solutions
 - Don't page.
 - Instead of swapping data to internal storage, Android and iOS start killing processes.
 - Add more RAM.
 - Use memory compression!
 - First version of Windows to support memory compression: Windows 10
 - First version of macOS to support memory compression: OS X Mavericks
 - zram, Linux kernel module for memory compression, released in 2014.

The Processor

- Memory Compression Pros
 - Faster than swapping
 - For systems with SSDs, it prevents premature death
- Memory Compression Cons
 - Compressing and decompressing memory contents takes CPU time
- Memory compression does not replace swapping. It supplements it.

The Processor

- x86 General Purpose 32-bit Registers

- EAX

- Extended Accumulator Register
- Used for I/O port access, arithmetic operations, interrupts, etc.

- EBX

- Extended Base Register
- Used as a base pointer for memory access

- ECX

- Extended Count Register
- Used as a loop counter

- EDX

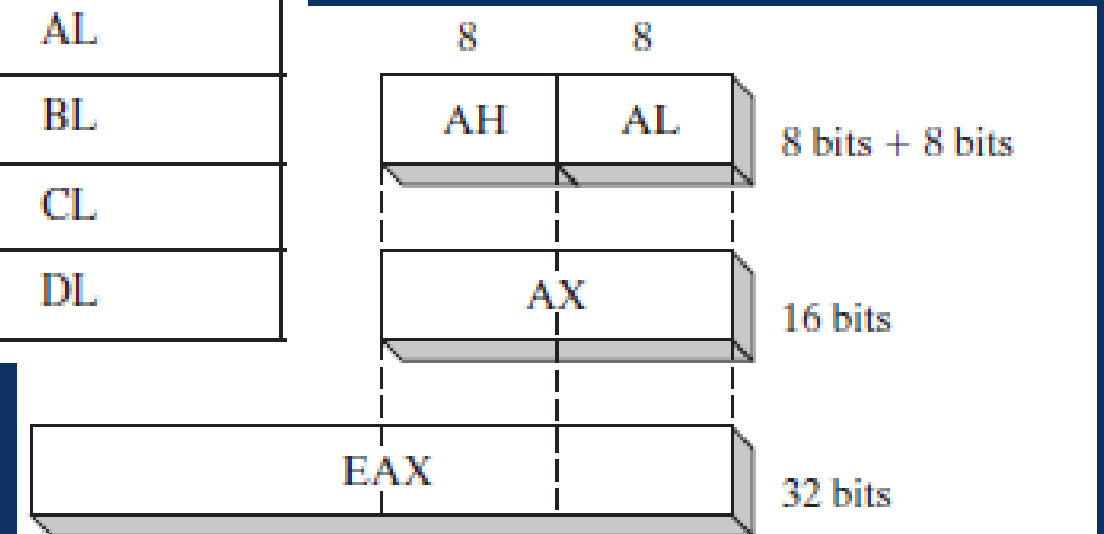
- Extended Data Register
- Used to store various data

In practice, these registers can be—and often are—used for storing whatever data the executing program needs.

The Processor

- x86 General Purpose 32-bit Registers
 - Some registers can be accessed in segments

32-Bit	16-Bit	8-Bit (High)	8-Bit (Low)
EAX	AX	AH	AL
EBX	BX	BH	BL
ECX	CX	CH	CL
EDX	DX	DH	DL



The Processor

- x86 General Purpose 32-bit Registers
 - ESP
 - Extended Stack Pointer Register
 - Used as the stack pointer
 - EBP
 - Extended Base Pointer Register
 - Used to reference the bottom of the stack
 - ESI
 - Extended Source Index
 - Used for high-speed memory transfer
 - EDI
 - Extended Data Index
 - Used for high-speed memory transfer

These registers can be used for storing arbitrary data as well, but it's considered bad practice.

The Processor

- More x86 Registers
 - Segment Registers
 - In protected mode, these hold pointers to segment descriptor tables.
 - Segment descriptor tables store information about where the process's segment information is stored.
 - Instruction Pointer
 - Stores the memory address of the next instruction waiting to be executed.
 - Certain instructions modify this value, causing the program to jump to instructions located elsewhere.

The Processor

Instruction Jump Example

```
int x = (some value);  
if (x == 4)  
{  
    x += 1;  
}  
else  
{  
    x -= 1;  
}  
...
```

Memory Location	Instruction
01	Assign (###) to 'x'
02	Compare 'x' with 4
03	If true, change EIP to 04, otherwise change EIP to 06
04	Add 1 to 'x'
05	Change EIP to 07
06	Subtract 1 from 'x'
07	(next instruction)
...	(rest of program)

The Processor

- Flag Registers
 - These registers used each of their bits to store a flag that either controls an operation or reflects the outcome of an operation.
 - 1 = set
 - 0 = clear/reset
- Flags
 - Control Flags
 - These control the CPU's operation
 - Can be used to break after each instruction, stop execution when an overflow occurs, etc.
 - Status Flags
 - These are set based on the outcome of an arithmetic or logical operation

The Processor

- Status Flags
 - Carry Flag: Set when the last arithmetic operation (unsigned) resulted in a carry bit that was beyond the size of the destination.

Unsigned 32-bit example:

$$\begin{array}{r} 1\ 1\ \dots\ 1\ 1\ 1\ 1\ 1\ 1 \\ +\ 0\ 0\ \dots\ 0\ 0\ 0\ 0\ 0\ 1 \\ \hline 0\ 0\ \dots\ 0\ 0\ 0\ 0\ 0\ 0 \end{array}$$

CF will be set to 1 to indicate the overflow.


The Processor

- Status Flags
 - Overflow Flag
 - Set when the last arithmetic operation (signed) resulted in a carry bit that was beyond the size of the destination.
 - Sign Flag
 - Set when the last arithmetic operation resulted in a negative value.
 - Zero Flag
 - Set when the last arithmetic operation resulted in a zero result.

The Processor

- Status Flags
 - Auxiliary Carry Flag
 - Set when the last arithmetic operation resulted in a carry bit from the low nibble (lowest four bits) to the high nibble (highest four bits).
 - Parity Flag
 - Set when the least significant byte results in an even number of 1 bits.
 - Used for data validation.

The Processor



Overflow
Flag

Sign
Flag

Zero
Flag

Auxiliary
Carry
Flag

Parity
Flag

Carry
Flag

The Processor

- Additional Registers – Used to improve performance for certain tasks.
 - MMX Registers (Multi-Media Extension)
 - A set of 8 64-bit registers that are useful for improving performance of multimedia and communications applications
 - Float Point Unit (FPU)
 - Registers used for high-speed floating point arithmetic.
 - XMM Registers
 - 128-bit registers for doing work that is performance-critical

Next Week

- The CPU: Part 3
 - Floating Point Units
 - A little more about x86 and amd64 processors
 - Components of a computer
 - Memory Types
 - Input/Output Systems
 - I/O Access Levels
- Homework #1 due on Canvas at the start of class.